

GROUP ASSIGNMENT

GROUP MEMBERS	REG NO
1. LINO NGOR MAJOK MANGOR	2405000678
2. YVES BAHATI	2305000895
3. EDGAR GAPFIZI BAHATI	25011938
4. ELYSÉE NIYIBIZI	2305000921
5. DANIEL NIZEYIMANA	2401001868

Q1.

Creating a pandas data frame

**Viewing the dataset,
summarizing contents, identifying invalid data,
converting data type columns,
and treating all missing values**

```
▶ import numpy as np
    import matplotlib.pyplot as plt

    import pandas as pd
    data= {
        'Student_ID': ['S1', 'S2', 'S3', 'S4', 'S5', 'S6', 'S7', 'S8', 'S9'],
        'Age': [20, 21, 22, np.nan, 19, 22, 'twenty', 23, 24, ],
        'Score': [85, np.nan, 78, 90, 88, 92, np.nan, 87, 80],
        'Hours_studied': [10, 15, 7, np.nan, 12, 9, 14, 11, 8]
    }
    df=pd.DataFrame(data)
    print(df)
    df.isnull().sum()

    df['Age']=pd.to_numeric(df['Age'], errors='coerce')
    df.info()

    df["Age"].fillna(df["Age"].median(), inplace=True)
    df["Score"].fillna(df["Score"].median(), inplace=True)
    df["Hours_studied"].fillna(df["Hours_studied"].median(), inplace=True)
    print(df)
    df.isnull().sum()
```

```
import pandas as pd
import matplotlib.pyplot as plt

x = [10, 15, 7, 10, 12, 9, 14, 11, 8]
y = [85, 87, 78, 90, 88, 92, 87, 87, 80]

plt.scatter(x, y)
plt.title("Student Score And Studied Hours")
plt.xlabel("Score")
plt.ylabel("Studied Hours")
plt.grid(True)
plt.show()
```


```
import pandas as pd
import matplotlib.pyplot as plt

x = [10, 15, 7, 10, 12, 9, 14, 11, 8]
y = [85, 87, 78, 90, 88, 92, 87, 87, 80]

plt.scatter(x, y)
plt.title("Student Score And Studied Hours")
plt.xlabel("Score")
plt.ylabel("Studied Hours")
plt.grid(True)
plt.show()
```

```
...     Student_ID      Age   Score  Hours_studied
... 0         S1        20    85.0       10.0
  1         S2        21     NaN       15.0
  2         S3        22    78.0        7.0
  3         S4        NaN    90.0       NaN
  4         S5        19    88.0       12.0
  5         S6        22    92.0        9.0
  6         S7    twenty     NaN       14.0
  7         S8        23    87.0       11.0
  8         S9        24    80.0        8.0
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 9 entries, 0 to 8
Data columns (total 4 columns):
 #   Column           Non-Null Count  Dtype  
---  --  
 0   Student_ID      9 non-null      object 
 1   Age              7 non-null      float64
 2   Score            7 non-null      float64
 3   Hours_studied   8 non-null      float64
dtypes: float64(3), object(1)
memory usage: 420.0+ bytes
```

	Student_ID	Age	Score	Hours_studied
0	S1	20.0	85.0	10.0
1	S2	21.0	87.0	15.0
2	S3	22.0	78.0	7.0
3	S4	22.0	90.0	10.5
4	S5	19.0	88.0	12.0
5	S6	22.0	92.0	9.0
6	S7	22.0	87.0	14.0
7	S8	23.0	87.0	11.0
8	S9	24.0	80.0	8.0



Q2.

A.

A cleaner Dataset with fewer missing values is more reliable for analysis and model training. Enhanced model performance: Properly handling missing values helps models perform better by training on complete data, leading to more accurate predictions.

```
[154]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
Copy this cell (C)
data = {
    'YearsExperience': [3, 5, np.nan, 2, 4, np.nan],
    'Monthlysalary': [450, 5200, 4800, np.nan, 5100, 5300],
    'Team': ['Development', 'Marketing', 'Development', 'Sales', 'Marketing', 'Sales'],
    'EducationLevel': ['Bachelor\'s', 'Master\'s', None, 'Bachelor\'s', 'Bachelor\'s', 'Master\'s'],
    'RemoteWorkStatus': ['Yes', 'No', 'Yes', 'Maybe', 'Yes', 'No']
}

df = pd.DataFrame(data)

print('----- DATASET -----')
print('\n')
print(df)
print('\n')
print('----- DATA INFORMATION -----')
print('\n')
#print('\n')
df.info()

print('\n')
print('----- MISSING VALUES -----')
print(df.isnull().sum())

----- DATASET -----
```

	YearsExperience	Monthlysalary	Team	EducationLevel	RemoteWorkStatus
0	3.0	450.0	Development	Bachelor's	Yes
1	5.0	5200.0	Marketing	Master's	No
2	NaN	4800.0	Development	None	Yes
3	2.0	NaN	Sales	Bachelor's	Maybe
4	4.0	5100.0	Marketing	Bachelor's	Yes
5	NaN	5300.0	Sales	Master's	No

```
----- DATA INFORMATION -----
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 6 entries, 0 to 5
Data columns (total 5 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   YearsExperience  4 non-null      float64
 1   Monthlysalary   5 non-null      float64
 2   Team             6 non-null      object  
 3   EducationLevel  5 non-null      object  
 4   RemoteWorkStatus 6 non-null      object  
dtypes: float64(2), object(3)
memory usage: 372.0+ bytes
```

```
----- MISSING VALUES -----
```

	YearsExperience	Monthlysalary	Team	EducationLevel	RemoteWorkStatus
	2	1	0	1	0

```
[160]: df['YearsExperience'] = df['YearsExperience'].fillna(df['YearsExperience'].mean())
df['Monthlysalary'] = df['Monthlysalary'].fillna(df['Monthlysalary'].median())
df['EducationLevel'] = df['EducationLevel'].fillna(df['EducationLevel'].mode()[0])

print('----- MISSING VALUES AFTER BEING TREATED -----')
print(df.isnull().sum())
```

```
----- MISSING VALUES AFTER BEING TREATED -----
YearsExperience    0
Monthlysalary      0
Team                0
EducationLevel      0
RemoteWorkStatus    0
dtype: int64
```

```
[162]: print('\n\n')
print('----- DATA DESCRIPTION -----')
print(df.describe())
```

```
----- DATA DESCRIPTION -----
   YearsExperience  Monthlysalary
count      6.000      6.000000
mean       3.500    4325.000000
std        1.000   1905.715089
min        2.000   450.000000
25%       3.125   4875.000000
50%       3.500   5100.000000
75%       3.875   5175.000000
max        5.000   5300.000000
```

B. CALCULATING MEDIAN FOR MONTHLY SALARY

```
[164]: print('----- CALCULATING MEDIAN FOR MONTHLY SALARY -----')

print('CHECK IF ALL VALUES ARE TREATED PROPERLY')
print(df['Monthlysalary'].isnull())

# IF YES CONTINUE TO CALCULATE THE MEDIAN USING THE METHOD
print('\n')
print(f'MONTHLY SALARY MEDIAN WILL BE >>>> {df['Monthlysalary'].median()}')

----- CALCULATING MEDIAN FOR MONTHLY SALARY -----
CHECK IF ALL VALUES ARE TREATED PROPERLY
0  False
1  False
2  False
3  False
4  False
5  False
Name: Monthlysalary, dtype: bool

MONTHLY SALARY MEDIAN WILL BE >>>> 5100.0
```

C.

I. GENERATING A NEW COLUMN FOR PERFORMANCE SCORE

```
[170]: # GENERATING A NEW COLUMN FOR PERFORMANCE SCORE
# FOR REPRODUCTIVITY
np.random.seed(42)
df['PerformanceScore'] = np.random.randint(50, 101, size=len(df))
print(df)
```

	YearsExperience	MonthlySalary	Team	EducationLevel
0	3.0	450.0	Development	Bachelor's
1	5.0	5200.0	Marketing	Master's
2	3.5	4800.0	Development	Bachelor's
3	2.0	5100.0	Sales	Bachelor's
4	4.0	5100.0	Marketing	Bachelor's
5	3.5	5300.0	Sales	Master's


	RemoteWorkStatus	PerformanceScore
0	Yes	88
1	No	78
2	Yes	64
3	Maybe	92
4	Yes	57
5	No	70

II. YEARS OF EXPERIENCE VERSUS PERFORMANCE SCORE

```
[172]: print('----- PLOTTING -----')
# Years of Experience versus Performance Score
# Get unique categories
statuses = df['RemoteWorkStatus'].unique()

plt.figure()
for status in statuses:
    subset = df[df['RemoteWorkStatus'] == status]
    plt.scatter(subset['YearsExperience'], subset['PerformanceScore'], label=status)

plt.scatter(df['YearsExperience'], df['PerformanceScore'])
plt.xlabel('Years of Experience')
plt.ylabel('Performance Score')
plt.title('Years of experience versus Performance score plot')
plt.legend()
plt.show()
```



Question 3.

A.

```
[24]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error

# Global variables
separator = " " * 20
divider = f"\n{'-'*80}\n"

# Given Dataset
data = {
    'YearsExperience': [1.1, 2.0, 3.2, 4.5, np.nan, 6.8, 7.5, 8.3, 9.0, 10.5],
    'Salary': [39343, 46205, np.nan, 60000, 65200, 72500, np.nan, 83000, 88000, 95000]
}

df = pd.DataFrame(data)

# Remove rows with missing values
df_clean = df.dropna()

x = df_clean[['YearsExperience']].values
y = df_clean['Salary'].values
print(f'{separator} Cleaned DataFrame: {separator}')
print(df_clean)

_____  
Cleaned DataFrame:  
_____  
YearsExperience Salary  
0 1.1 39343.0  
1 2.0 46205.0  
3 4.5 60000.0  
5 6.8 72500.0  
7 8.3 83000.0  
8 9.0 88000.0  
9 10.5 95000.0
```

```
[17]: # Split into training and testing sets
x_train, x_test, y_train, y_test = train_test_split(
    x, y, test_size=0.2, random_state=42
)

print(f'{separator} x_train: {separator}')
print(x_train)
print(divider)
print(f'{separator} y_train: {separator}')
print(y_train)

_____  
x_train:  
[[ 9. ]  
[ 4.5]  
[ 8.3]  
[ 6.8]  
[10.5]]
```

```
y_train: [88000. 60000. 83000. 72500. 95000.]  
[20]: # Training the simple Linear regression model  
regressor = LinearRegression()  
regressor.fit(x_train, y_train)  
  
print(f"{separator} Model trained! {separator}")  
Model trained!  
[21]: # Viewing Learned parameters  
print(f"{separator} Slope (Coefficient): {separator}")  
print(regressor.coef_[0])  
print(f"{separator} Intercept: {separator}")  
print(regressor.intercept_)  
Slope (Coefficient): 5996.262219666475  
Intercept: 32809.22944220816
```

B.

```
[30]: # Predicting on the test set
y_pred = regressor.predict(x_test)

print(f"{separator} Predicted values: {separator}")
print(y_pred)
print(divider)
print(f"{separator} Actual values: ")
print(y_test)

----- Predicted values: -----
[39405.11788384 44801.75388154]

-----
----- Actual values:
[39343. 46205.]
```

```
[31]: # Computing Mean Squared Error
mse = mean_squared_error(y_test, y_pred)
print(f"{separator} Mean Squared Error: {separator}")
print(mse)

----- Mean Squared Error:
986479.1502314303
```

```
[32]: # Computing Root Mean Squared Error for easier interpretation
rmse = np.sqrt(mse)
print(f"{separator} Root Mean Squared Error: {separator}")
print(rmse)

----- Root Mean Squared Error:
993.216567638413
```

C.

```
[33]: # Extracting and printing slope and intercept
slope = regressor.coef_[0]
intercept = regressor.intercept_

print(f"{separator} Slope (Coefficient): {separator}")
print(slope)
print(divider)
print(f"{separator} Intercept: {separator}")
print(intercept)

_____  
Slope (Coefficient): _____  
5996.262219666475  
-----  
_____  
Intercept: _____  
32809.22944220816
```

```
[36]: # Interpreting the coefficient
if slope > 0:
    print(">> Positive coefficient: Salary increases as YearsExperience increases.")
elif slope < 0:
    print(">> Negative coefficient: Salary decreases as YearsExperience increases.")
else:
    print(">> Zero coefficient: No relationship between YearsExperience and Salary.")

print(divider)
print(f"{{separator}} Interpretation: {{separator}}")
print(f"{{>>}} For each additional year of experience, salary increases by about {slope:.2f}.")
print(f"{{>>}} When experience is zero, predicted salary is about {intercept:.2f}.")

>> Positive coefficient: Salary increases as YearsExperience increases.
-----
Interpretation:
>> For each additional year of experience, salary increases by about 5996.26.
>> When experience is zero, predicted salary is about 32809.23.
```

[]:

